



# CLOUD COMPUTING CONCEPTS

with Indranil Gupta (Indy)

## P2P SYSTEMS

Lecture E

---

### CHORD

# DHT=DISTRIBUTED HASH TABLE

- A hash table allows you to insert, lookup, and delete objects with keys
- A *distributed* hash table allows you to do the same in a distributed setting (objects=files)
- Performance concerns:
  - Load balancing
  - Fault-tolerance
  - Efficiency of lookups and inserts
  - Locality
- Napster, Gnutella, FastTrack are all DHTs (sort of)
- So is Chord, a structured peer-to-peer system that we study next

# COMPARATIVE PERFORMANCE

	Memory	Lookup Latency	#Messages for a lookup	
Napster	$O(1)$ ( $O(N)$ @server)	$O(1)$	$O(1)$	
Gnutella	$O(N)$	$O(N)$	$O(N)$	

# COMPARATIVE PERFORMANCE

	Memory	Lookup Latency	#Messages for a lookup	
Napster	$O(1)$ ( $O(N)$ @server)	$O(1)$	$O(1)$	
Gnutella	$O(N)$	$O(N)$	$O(N)$	
Chord	$O(\log(N))$	$O(\log(N))$	$O(\log(N))$	

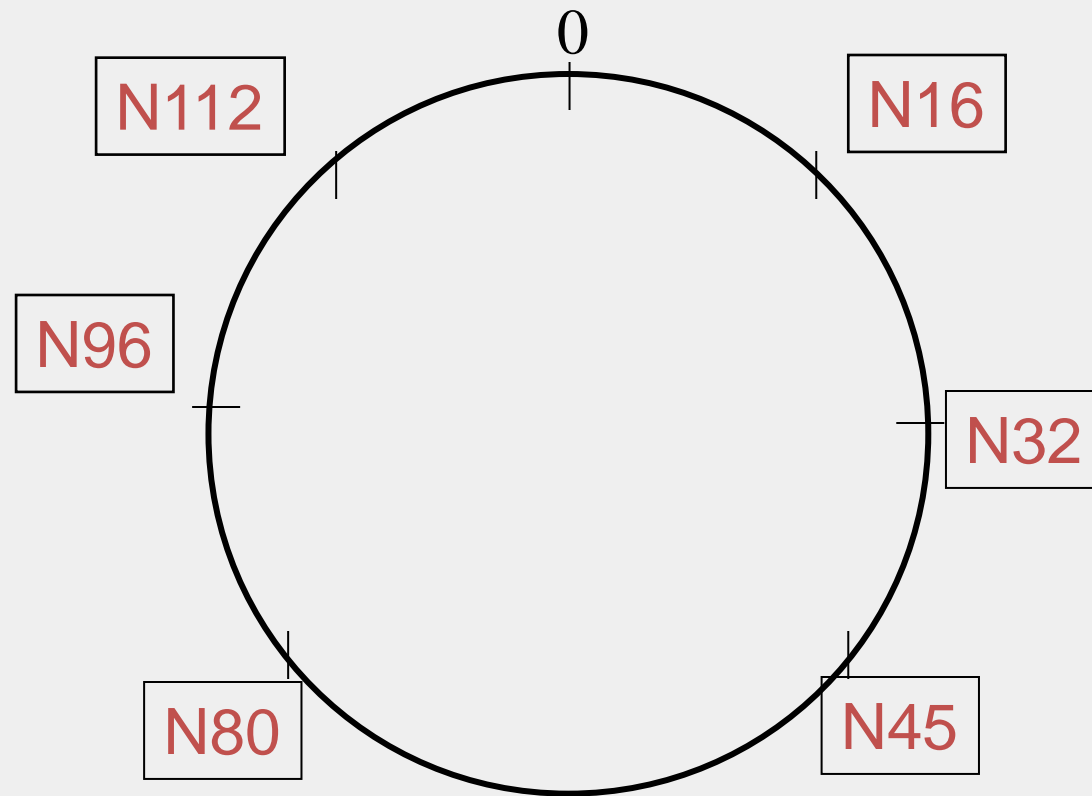
# CHORD

- Developers: I. Stoica, D. Karger, F. Kaashoek, H. Balakrishnan, R. Morris, Berkeley, and MIT
- Intelligent choice of neighbors to reduce latency and message cost of routing (lookups/inserts)
- Uses *Consistent Hashing* on node's (peer's) address
  - **SHA-1**(ip\_address,port) → 160 bit string
  - Truncated to  $m$  bits
  - Called peer *id* (number between 0 and  $2^m - 1$  )
  - Not unique but id conflicts very unlikely
  - Can then map peers to one of  $2^m$  logical points on a circle

# RING OF PEERS

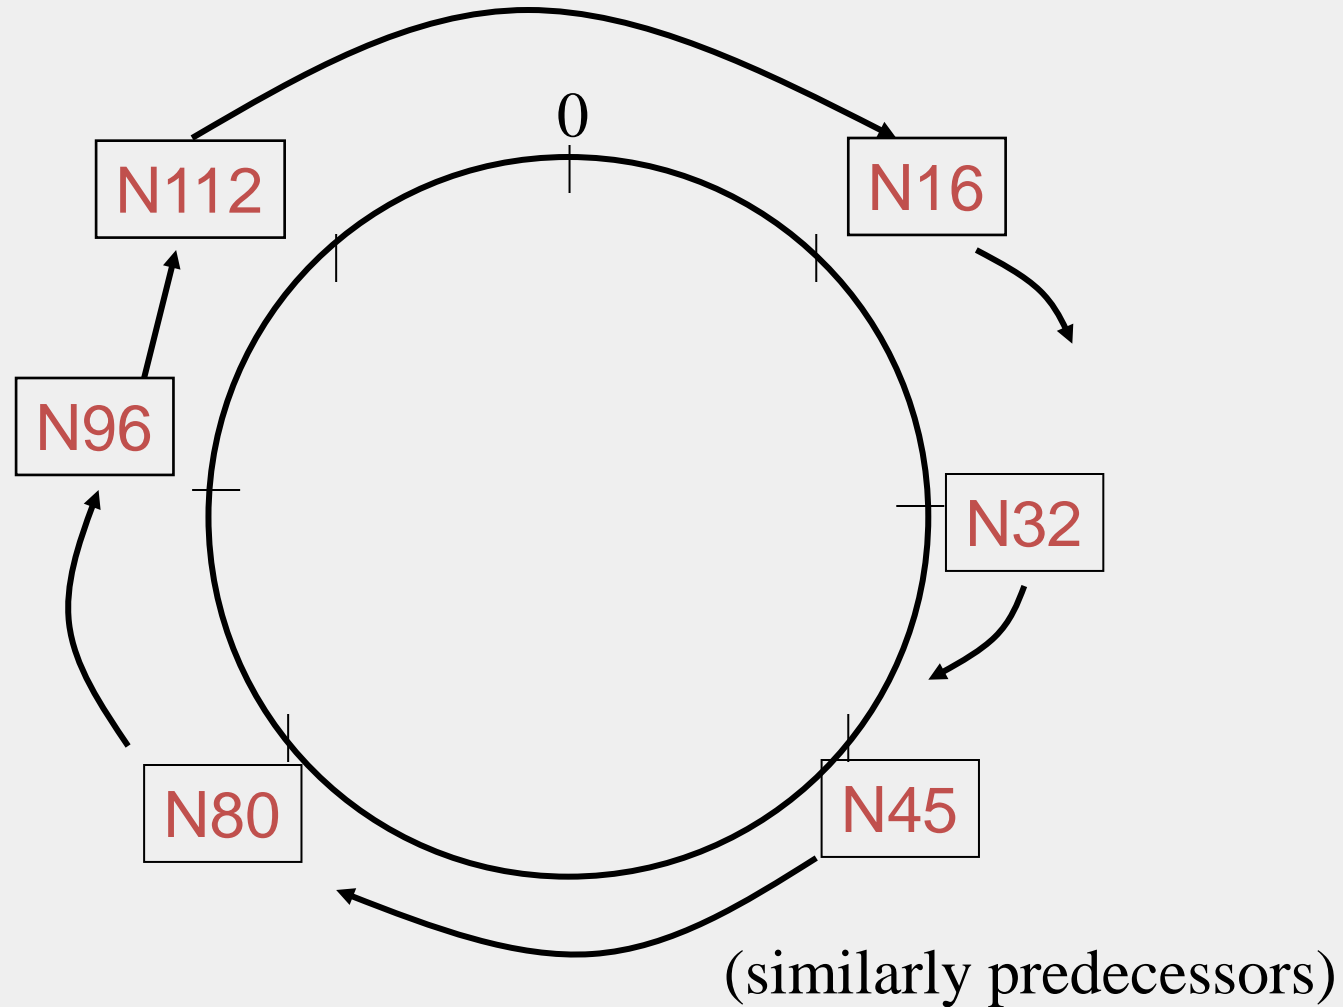
Say  $m=7$

6 nodes



# PEER POINTERS (1): SUCCESSORS

Say  $m=7$




# PEER POINTERS (2): FINGER TABLES

Say  $m=7$

Finger Table at N80

$i$	$ft[i]$
0	96
1	96
2	96
3	96
4	96
5	112
6	16



$i$ th entry at peer with id  $n$  is first peer with id  $\geq n + 2^i \pmod{2^m}$

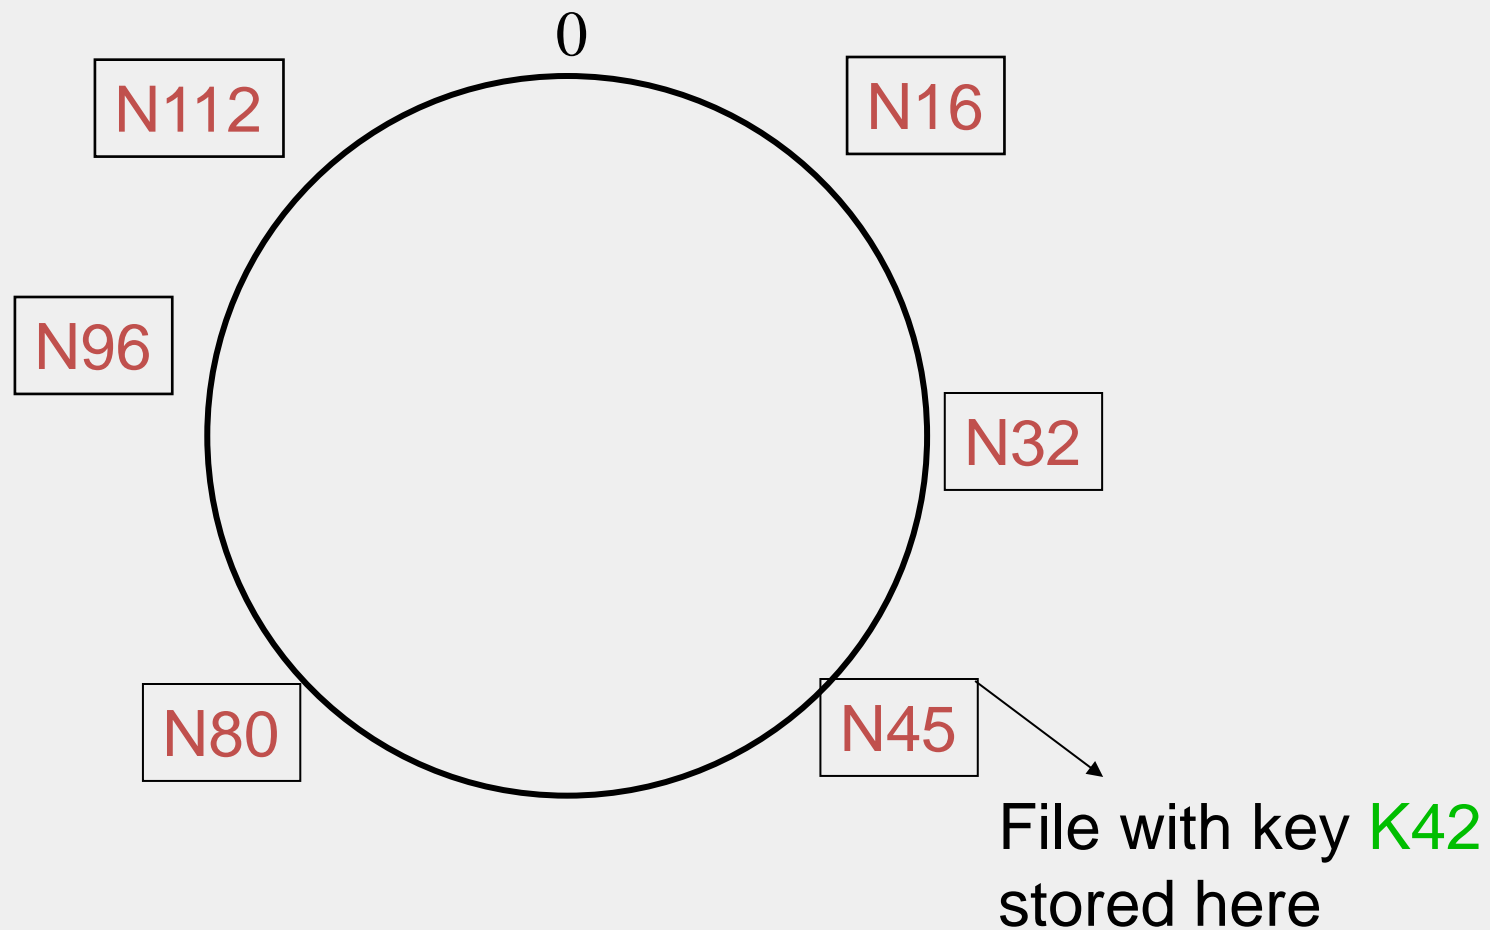


# WHAT ABOUT THE FILES?

- Filenames also mapped using same consistent hash function
  - $\text{SHA-1}(\text{filename}) \rightarrow 160 \text{ bit string (key)}$
  - File is stored at **first peer with id greater than its key (mod  $2^m$ )**
- File *cnn.com/index.html* that maps to key K42 is stored at first peer with id greater than 42
  - Note that we are considering a different file-sharing application here: *cooperative web caching*
  - The same discussion applies to any other file sharing application, including that of mp3 files.
- Consistent Hashing  $\Rightarrow$  with K keys and N peers, each peer stores  $O(K/N)$  keys. (i.e.,  $< c.K/N$ , for some constant  $c$ )

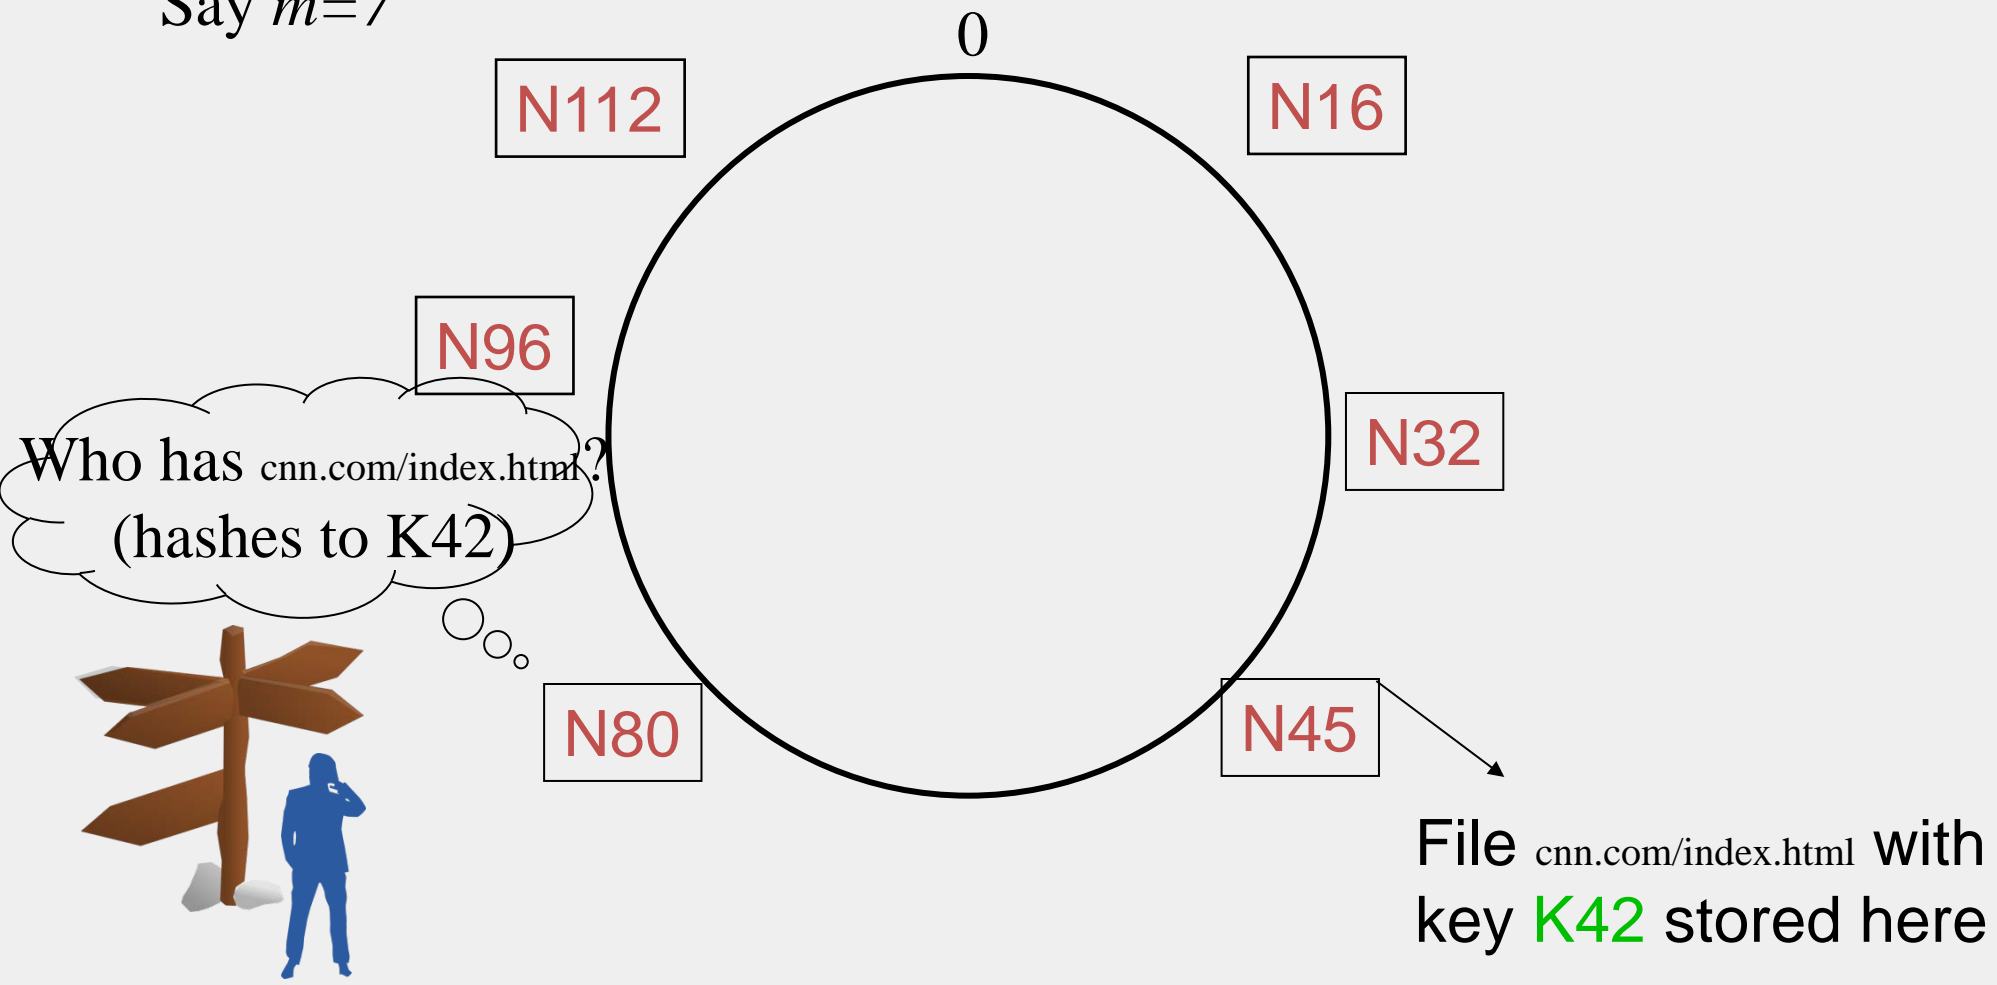
# MAPPING FILES

Say  $m=7$



# SEARCH

Say  $m=7$



# SEARCH

At node  $n$ , send query for key  $k$  to largest successor/finger entry  $\leq k$   
if none exist, send query to  $successor(n)$

0

N112

N16

Say  $m=7$

N96

N32

Who has `cnn.com/index.html`?  
(hashes to K42)

N80

N45




File `cnn.com/index.html` with  
key **K42** stored here

# SEARCH

At node  $n$ , send query for key  $k$  to largest successor/finger entry  $\leq k$   
if none exist, send query to  $successor(n)$

Say  $m=7$



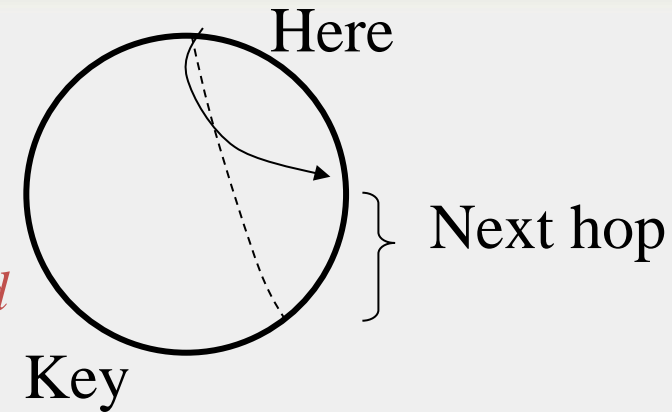
# ANALYSIS

Search takes  $O(\log(N))$  time

## Proof

- (Intuition): *at each step, distance between query and peer-with-file reduces by a factor of at least 2*
- (Intuition): after  $\log(N)$  forwardings, distance to key is at most  $2^m / 2^{\log(N)} = 2^m / N$
- Number of node identifiers in a range of  $2^m / N$  is  $O(\log(N))$  with high probability (why? SHA-1! and “Balls and Bins”)

So using *successors* in that range will be ok, using another  $O(\log(N))$  hops



# ANALYSIS (CONTD.)

- $O(\log(N))$  search time holds for file insertions too (in general for *routing to any key*)
  - “Routing” can thus be used as a **building block** for
    - All operations: insert, lookup, delete
- $O(\log(N))$  time true only if finger and successor entries correct
- When might these entries be wrong?
  - When you have failures